

R Examples from BIOST 511 - Set 7 and 8

Taught by Lurdes Inoue, Ph.D.

Brian D. Williamson

University of Washington
Department of Biostatistics

November 14, 2014

Disclaimer

All of the following slides and functions are to be used when given sample descriptive statistics

For raw data, refer to the documentation for `ttest()` and other functions given on <http://www.emersonstatistics.com/R>

Test of Proportions

Use the `prop.test()` function from the `stats` package
(automatically loaded into R along with the base package)

```
prop.test(110, 200, p=.3, alternative="two.sided")
```

1-sample proportions test with continuity correction

```
data: 110 out of 200, null probability 0.3  
X-squared = 58.3393, df = 1, p-value = 2.206e-14  
alternative hypothesis: true p is not equal to 0.3  
95 percent confidence interval:  
 0.4782703 0.6197775  
sample estimates:  
      p  
0.55
```

– If we want a different test, we change the `alternative` argument

t-test using Sample Descriptive Statistics

Use the `ttesti()` function from the `uwIntroStats` package

```
ttesti(24, 175, 35, 43, null.hyp=230)
```

One-sample t-test

	Obs	Mean	Std. Error	SD	95 %CI
x	24	175	7.1443	35	160.221 189.779

t-statistic = -7.698396 , df = 23

Ho: mean = 230

Ha: mean != 230 , $\Pr(|T| > |t|) = 8.24e-08$

– Again, if we want a different test, we change the `alternative` argument

Binomial Test

Use the `binom.test()` function from the `stats` package

```
binom.test(110, 200, p=.3, alternative="two.sided")
```

Exact binomial test

```
data: 110 and 200
number of successes = 110, number of trials = 200, p-value = 3.125e-13
alternative hypothesis: true probability of success is not equal to 0.3
95 percent confidence interval:
 0.4782493 0.6202464
sample estimates:
probability of success
      0.55
```

– Again, if we want a different test, we change the `alternative` argument

Group 1 and Group 2 One-Sample Tests

```
## One sample t-test for Group 1  
ttesti(obs=10, mean=-1.6, sd=1.5, null.hyp=0)
```

```
One-sample t-test  
  Obs Mean Std. Error SD  95 %CI  
x  10  -1.6  0.4743    1.5 -2.673 -0.527
```

```
t-statistic = -3.373096 , df = 9
```

```
Ho: mean = 0
```

```
Ha: mean != 0 , Pr(|T| > |t|) = 0.0082165
```

```
## For Group 2  
ttesti(obs=30, mean=-.7, sd=2.1, null.hyp=0)
```

```
One-sample t-test  
  Obs Mean Std. Error SD  95 %CI  
x  30  -0.7  0.3834    2.1 -1.484 0.084
```

```
t-statistic = -1.825742 , df = 29
```

```
Ho: mean = 0
```

```
Ha: mean != 0 , Pr(|T| > |t|) = 0.078203
```

Difference of Means with Sample Descriptive Statistics

Unequal Variances

Use the `ttesti()` function in the `uwIntroStats` package
For the Group 1 Group 2 example,

```
## Two sample for the difference of means, with variance unequal  
ttesti(10, -1.6, 1.5, 30, -.7, 2.1)
```

```
Two-sample t-test with unequal variances  
  Obs Mean Std. Error SD  95 %CI  
x   10  -1.6  0.4743    1.5 -2.673 -0.527  
y   30  -0.7  0.3834    2.1 -1.484  0.084  
diff 40  -0.9  0.6099    NA  -2.135  0.335
```

```
t-statistic = -1.475608 , Satterthwaite's df = 21.72386
```

```
Ho: mean(x) - mean(y) = diff = 0
```

```
Ha: diff != 0 , Pr(|T| > |t|) = 0.1544
```

Proportions

Use the `binom.test()` function in the stats package

```
## Exact binomial test  
binom.test(7, 26, .5)
```

Exact binomial test

```
data: 7 and 26  
number of successes = 7, number of trials = 26, p-value = 0.02896  
alternative hypothesis: true probability of success is not equal to 0.5  
95 percent confidence interval:  
 0.1157322 0.4778748  
sample estimates:  
probability of success  
 0.2692308
```


Chi-squared Test

Use the `chisq.test()` function in the `stats` package
If given raw data, use the `tabulate()` function in the
`uwIntroStats` package

The `correct=FALSE` argument tells R to NOT use Yates' continuity correction

For this small data example, you will get a warning from R, because the expected cell counts are too small for a reliable Chi-squared test

```
chisq.test(matrix(c(2,23,5,30), nrow=2, byrow=T), correct=FALSE)
```

Pearson's Chi-squared test

```
data:  matrix(c(2, 23, 5, 30), nrow = 2, byrow = T)
X-squared = 0.5591, df = 1, p-value = 0.4546
```

Fisher's Exact Test

Use the `fisher.test()` function in the `stats` package
If given raw data, use the `tabulate()` function in the
`uwIntroStats` package with the correct arguments

```
fisher.test(matrix(c(2,23,5,30), nrow=2, byrow=T))
```

Fisher's Exact Test for Count Data

```
data:  matrix(c(2, 23, 5, 30), nrow = 2, byrow = T)
p-value = 0.6882
alternative hypothesis: true odds ratio is not equal to 1
95 percent confidence interval:
 0.04625243 3.58478157
sample estimates:
odds ratio
 0.527113
```

Chi-squared Goodness of Fit Test

Use the `chisq.test()` function in the `stats` package
Simply enter in a vector rather than a matrix and a goodness of fit test will be performed

Homogeneity Test

Use the `chisq.test()` function in the `stats` package
If given raw data, use the `tabulate()` function in the
`uwIntroStats` package with the correct arguments

```
## Chi-squared example  
chisq.test(matrix(c(7, 55, 489, 475,  
                   293, 38, 61, 129, 570, 431, 154, 12),  
                 nrow=2, byrow=T))
```

Pearson's Chi-squared test

```
data: matrix(c(7, 55, 489, 475, 293, 38, 61, 129, 570, 431, 154, 12), nrow = 2, byrow = T)  
X-squared = 137.7193, df = 5, p-value < 2.2e-16
```

Independence Test

```
chisq.test(matrix(c(52, 79, 342, 226, 62, 153, 417, 262, 53, 213, 629, 375,  
                    54, 231, 571, 244, 18, 46, 139, 74, 25, 139, 330, 116),  
                  nrow=4, byrow=T))
```

Pearson's Chi-squared test

```
data: matrix(c(52, 79, 342, 226, 62, 153, 417, 262, 53,  
213, 629, 375, 54, 231, 571, 244, 18, 46, 139, 74, 25, 139, 330, 116), nrow = 4, byrow = T)  
X-squared = 2041.811, df = 15, p-value < 2.2e-16
```

Sign Test

No direct analog for this test in R

Since it ignores a lot of information, use the binomial test or McNemar's test

Wilcoxon Signed Rank Test

Use the `wilcoxon` function from the `uwIntroStats` package
First you must have data vectors, so create them if necessary

```
cf <- c(1153, 1132, 1165, 1460, 1162, 1493, 1358, 1453, 1185, 1824, 1793, 1930, 2075)
healthy <- c(996, 1080, 1182, 1452, 1634, 1619, 1140, 1123, 1113, 1463, 1632, 1614, 1836)
wilcoxon(cf, healthy, paired=TRUE)
```

```
Wilcoxon signed rank test
      obs sum ranks expected
positive  10      71      45.5
negative   3      20      45.5
zero       0       0       0.0
all       13      91     91.0
```

```
unadjusted variance  204.75
adjustment for ties   0.00
adjustment for zeroes 0.00
adjusted variance    204.75
      HO Ha
Hypothesized Median 0 two.sided
Test Statistic p-value
V 71             0.080322
Z 1.7821          0.037368
```

Wilcoxon Rank Sum Test

Use the `wilcoxon` function from the `uwIntroStats` package
First you must have data vectors, so create them if necessary

```
ebv <- c(2.9, 12.1, 2.6, 2.5, 2.8, 15.8, 3.2, 1.8, 7.8,
        2.9, 3.2, 8.0, 1.5, 6.3, 1.2, 3.5, 4.5, 1.3,
        1.0, 1.0, 1.3, 1.9, 1.3, 2.1, 2.1, 1.0)
sero <- c(rep(1,16), rep(0,10))
data <- cbind(ebv, sero)
wilcoxon(data[sero==1], data[sero==0])
```

```
Wilcoxon rank sum test
      obs rank sum expected
Y      20      369      689
X      32     1009      689
combined 52     1378     1378
```

```
unadjusted variance 2826.666667
adjustment for ties  2.968326
adjusted variance   2823.698341
      HO Ha
Hypothesized Median 0 two.sided
Test Statistic p-value
W 481          0.0018285
Z 3.1168       0.00091424
Warning message:
In wilcoxon.do(x = y, y = x, alternative = alternative, mu = mu, :
cannot compute exact p-value with ties
```

The warning message means that there are tied values, so the
p-values are approximate

Systolic Blood Pressure Example

Create the data set by saving it as a text file, then load it into R and attach it

```
systolic <- read.table("P:/TA/systolic.txt", header=TRUE, quote="\"", stringsAsFactors=FALSE)
```

Next regress systolic (the y-variable) vs age using the `regress()` function, with first argument "mean"

```
regress("mean", systolic$systolic, model=uModel(systolic$age))
```

```
Call:
regress(fnctl = "mean", y = systolic$systolic, model = uModel(systolic$age))
```

```
Residuals:
    Min       1Q   Median       3Q      Max
-7.1395 -2.3314 -0.2163  2.1872  5.8605
```

```
Coefficients:
            Estimate Naive SE Robust SE    95%L    95%H      F stat    df Pr(>F)
Intercept      67.68    3.191    2.440    62.45    72.91    769.25    1 < 0.00005
systolic$age     6.153    0.9283    0.6511    4.757    7.550    89.32    1 < 0.00005
```

```
Residual standard error: 3.403 on 14 degrees of freedom
```

```
Multiple R-squared:  0.7584, Adjusted R-squared:  0.7411
```

```
F-statistic: 43.94 on 1 and 14 DF,  p-value: 1.135e-05
```

The `uModel()` function gives the correct names and allows multiple variables to be listed in the `model` argument

Prediction of the mean

Again use the `regress()` function

We will need items from the list it returns, so get the predicted \hat{y} 's with `$linearPredictor` and the standard deviation with `$sigma`:

```
yhat <- sysRegress$linearPredictor
## get standard error about line
se <- sysRegress$sigma/sqrt(14)
## calculate confidence interval around the line using the quantile of the t distribution
lower <- yhat - qt(.975, 14)*se
upper <- yhat + qt(.975, 14)*se
cbind(round(lower, 3), round(yhat, 3), round(upper,3))
```

	[,1]	[,2]	[,3]
[1,]	84.189	86.140	88.090
[2,]	90.342	92.293	94.244
[3,]	84.189	86.140	88.090
[4,]	78.035	79.986	81.937
[5,]	90.342	92.293	94.244
[6,]	96.496	98.447	100.397
[7,]	78.035	79.986	81.937
[8,]	84.189	86.140	88.090
[9,]	96.496	98.447	100.397
[10,]	90.342	92.293	94.244
[11,]	78.035	79.986	81.937
[12,]	84.189	86.140	88.090
[13,]	84.189	86.140	88.090
[14,]	90.342	92.293	94.244
[15,]	84.189	86.140	88.090
[16,]	84.189	86.140	88.090

Prediction of a new value

Again use the regress() function

```
yhat <- sysRegress$linearPredictor
## get standard error about line
se <- sysRegress$sigma*sqrt(1+1/16+(systolic$age-mean(systolic$age))^2/(sd(systolic$age)^2))
## calculate confidence interval around the line using the quantile of the t distribution
lower <- yhat - qt(.975, 14)*se
upper <- yhat + qt(.975, 14)*se
cbind(round(lower, 3), round(yhat, 3), round(upper,3))
```

	[,1]	[,2]	[,3]
[1,]	78.240	86.140	94.039
[2,]	83.090	92.293	101.497
[3,]	78.240	86.140	94.039
[4,]	67.375	79.986	92.597
[5,]	83.090	92.293	101.497
[6,]	83.416	98.447	113.478
[7,]	67.375	79.986	92.597
[8,]	78.240	86.140	94.039
[9,]	83.416	98.447	113.478
[10,]	83.090	92.293	101.497
[11,]	67.375	79.986	92.597
[12,]	78.240	86.140	94.039
[13,]	78.240	86.140	94.039
[14,]	83.090	92.293	101.497
[15,]	78.240	86.140	94.039
[16,]	78.240	86.140	94.039